

## Table of Content

|                                  |    |
|----------------------------------|----|
| Architecture.....                | 2  |
| 1    Overview.....               | 2  |
| 2    Register .....              | 3  |
| 3    ISA Summary .....           | 3  |
| 1.1    3 Operand operations..... | 3  |
| 1.2    Immediate mode.....       | 4  |
| 4    Flags & Branch.....         | 5  |
| 5    ISA detail .....            | 6  |
| 6    Pseudo Instructions.....    | 21 |
| 7    Memory Map .....            | 21 |
| Micro-architecture .....         | 22 |
| 1    Overview.....               | 22 |
| 2    Memory Controller .....     | 23 |
| 3    Command Patcher.....        | 25 |
| 4    Patch Dispatcher.....       | 27 |
| 5    RT core.....                | 30 |
| 6    IC buffer.....              | 32 |
| 7    IC core .....               | 34 |
| 8    IP cores.....               | 36 |
| 9    Memory.....                 | 37 |
| 9.1    Memory overview.....      | 37 |
| 9.2    Instruction memory .....  | 37 |
| 9.3    Constant memory.....      | 38 |
| 9.4    Main memory.....          | 39 |
| 9.5    Triangle memory.....      | 40 |
| 10    Coding standard.....       | 42 |

# Architecture

## 1 Overview

Our processing unit is a parallel-processing unit that mimics the behavior of current GPU. From a software perspective, the same instructions will be executed for each pixel. Each thread will have a pixel ID showing its location on the final image. It will also be assigned a thread ID, which is between 31-0. This is the ID within each patch. The hardware will keep track of the pixel ID of each thread and place it in the register file before the thread starts. Each thread will use this ID to compute which pixel this thread is working on and find the correct camera ray. The PC and the stack-base pointer will also be prepared by the hardware and placed in register file.




Figure 1 Relationship between Thread ID and Pixel ID

Each thread can fully utilize any resources on its RT core. It has its own PC (we are no longer doing lock step), register file, integer compute unit, floating point compute unit, as well as memory access. We have two types of memory: constant memory, and local memory. Constant memory defines the global state, like camera location, direction, material information. In Computer Graphics terminology, this is the “Uniform”. All constant memory address is the same for each thread. The other type of memory is local memory, and it is used for stack and output. Each thread has its own memory space. Software can access these memories freely as there is only one thread, (since we support multiple reads performed at the same time). Our hardware will make sure this thread is accessing its own memory space. Constant memory is four-byte addressable, and it only supports scalar read. Main Memory is four-byte addressable and support scalar vector read write. Both memories use little endian formatting.

To access IC core accelerator, a thread evokes a special instruction called “trace”. Before calling trace, it needs to put the function call parameters in special locations in vector registers. When this instruction is executed, it will halt the thread and perform a hardware context switch. Hardware will store the current PC

and the stack pointer of the thread, and then will read from those vector registers. Then hardware will give this information to the IC core and schedule a different thread to this core. When the IC core is finished and hardware is ready to schedule this thread back, it will put the trace result, PC and stack pointer back to the register file before it executes. If software needs to preserve other registers, these registers must be pushed to the stack. Finally, each thread can return one vector back to the host as the pixel color. They need to write the return value to the last vector of their main memory space.

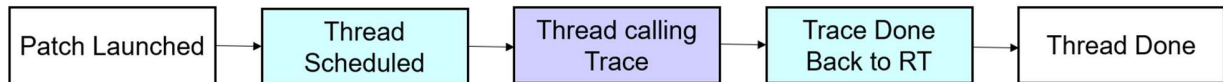


Figure 2 Typical Thread execution cycle

After careful deliberations, we decided to remove command processor and replace it with a state machine. Therefore, its behavior is fixed and only the RT core needs instructions.

## 2 Register

Each thread has 32 32-bit scalar registers from R0-R31. The R31 register is reserved for PC register. R0 is reserved for the value zero and it is read-only. R30 is reserved for link register. R29 is reserved for stack pointer and R28 is reserved for stack base pointer. We also provide 16 4x32-bit vector registers, RV0-RV15. Similarly, R0 is reserved for the value zero and it is read-only.

## 3 ISA Summary

We have specially designed the ISA to reduce the number of multiplexors in our decode process.

### 1.1 3 Operand operations.

All three operand operations do not have an immediate value and have a 0 for their first opcode bit.

Next 3 bits: Operand type

000: Vector \* Vector -> Vector

010: Vector \* Scalar -> Vector

011: Scalar \* Scalar -> Scalar (And Or Xor Not)

110: Scalar \* Scalar -> Scalar (Int)

111: Scalar \* Scalar -> Scalar (Int) Update Flag

100: Scalar \* Scalar -> Scalar (Float)

101: Scalar \* Scalar -> Scalar (Float) Update Flag

Next 2 bits: operation

00: Add (And)

01: Sub (Or)

10: Mul (Xor)

11: Div (Not)

| Opcode [31: 26] | Instruction | Explanation                                   |
|-----------------|-------------|---|
| 000000          | vv_add      | Add two vectors (vec4_float)                  |
| 000001          | vv_sub      | Subtract two vectors (vec4_float)             |
| 000010          | vv_mul_ele  | Elementwise multiply two vectors (vec4_float) |

|        |        |   |
|--------|--------|---|
| 000011 | vv_div | Divide two vectors (vec4_float)                         |
| 001000 | vf_add | Add a scaler/float to a vector (vec4_float)             |
| 001001 | vf_sub | Subtract a scaler/float from a vector (vec4_float)      |
| 001010 | vf_mul | Multiply a scaler/float to a vector (vec4_float)        |
| 001011 | vf_div | Divide a scaler/float from a vector (vec4_float)        |
| 001100 | not    | Bitwise not   |
| 001101 | and    | Bitwise and   |
| 001110 | or     | Bitwise or  |
| 001111 | xor    | Bitwise xor   |
| 011000 | ii_add | Add two integers  |
| 011001 | ii_sub | Subtract an integer by another integer                  |
| 011010 | ii_mul | Multiply two integers                                   |
| 011011 | ii_div | Divide an integer by another integer                    |
| 011101 | cmp_i  | Compare two integer scalars and output to flag register |
| 010000 | ff_add | Add two floats  |
| 010001 | ff_sub | Subtract a float by another float                       |
| 010010 | ff_mul | Multiply two floats                                     |
| 010011 | ff_div | Divide a float by another float                         |
| 010101 | cmp_f  | Compare two float scalars and output to flag register   |

## 1.2 Immediate mode

The last 16 bits of these instructions are either immediate values or “don’t care” values. These opcodes start with 1.

Next bits 0: This instruction has writeback.

Next bits 00: Integer core special function

Next bits 01: floating (vector) core special function

Next bits 10: Integer core operation with immediate

Next bits 11: memory load

Next bits 1: This instruction does not have writeback.

Next bits 00: Jump and branch.

Next bits 10: Program control

Next bits 11: memory store

| Opcode[31: 26] | Instruction  | Explanation   |
|----------------|--------------|---|
| 100000         | s_write_high | Write immediate into higher 16 bits of the register |
| 100001         | s_write_low  | Write immediate into lower 16 bits of the register  |
| 100010         | s_itof       | Convert integer to float                            |
| 100011         | s_ftoi       | Convert float to integer                            |
|                |              |   |
| 100100         | v_reduce     | Add all elements together                           |

|        |                |   |
|--------|----------------|---|
| 100101 | s_get_from_v   | Read one scalar from a vector into a register         |
| 100110 | s_sqrt         | Calculate the square root of the scalar (float)       |
| 100111 | v_get_from_s   | Put one scalar into a location of the vector register |
|        |                |   |
| 101000 | ii_addi        | Add an integer immediate to an integer                |
| 101001 | ii_subi        | Subtract an integer by an integer immediate           |
| 101010 | ii_muli        | Multiply an integer immediate and an integer          |
| 101011 | ii_divi        | Divide an integer by an integer immediate             |
|        |                |   |
| 101100 | s_load_4byte   | Load four bytes of scalar data from memory            |
| 101111 | v_load_16byte  | Load vector (four four-byte dimension) from memory    |
|        |                |   |
| 111100 | s_store_4byte  | Store four bytes of scalar data into memory           |
| 111111 | v_store_16byte | Store vector (four four-byte dimension) into memory   |
|        |                |   |
| 111010 | Trace          | Start a trace command                                 |
| 111011 | Fin            | Program ends  |
|        |                |   |
| 110000 | be             | Branch if equal                                       |
| 110001 | bne            | Branch if not equal                                   |
| 110010 | bl             | Branch If less than                                   |
| 110011 | bge            | Branch if greater or equal                            |
| 110100 | bg             | Branch if greater than                                |
| 110101 | ble            | Branch if less or equal                               |
| 110110 | jmp            | Jump to a label                                       |
| 110111 | ret            | Return to a stored address                            |

#### 4 Flags & Branch

Flag status is stored in a 2-bit register during the execution stage by only the “compare” instructions (cmp\_i, cmp\_f). Therefore, software needs to compare first and then evoke the branch instruction. Software developers are advised to avoid checking the zero flag after the cmp\_f command. This is because there might be some rounding error due to floating points. We advise software developers to use the following equation to check if two floats are the same.

$$|Op_1 - Op_2| < \varepsilon, \varepsilon = 0.001$$

| Value | Abbreviation | Description |
|-------|--------------|-------------|
| 00    | Z            | Zero        |
| 11    | N            | Negative    |
| 10    | P            | Positive    |

## 5 ISA detail

### 5.1 vv\_add

#### Syntax:

vv\_add <RVdst>, <RVa>, <RVb>

#### Pseudo Code: (Can be in C/Verilog)

$RVdst = RVa + RVb$

#### Encoding:

000000 xdddd xaaaa xbbbb xxx xxxx xxxx

#### Usage and Examples:

Add two 4x4byte vectors

vv\_add RV3, RV1, RV2

### 5.2 vv\_sub

#### Syntax:

vv\_sub <RVdst>, <RVa>, <RVb>

#### Pseudo Code: (Can be in C/Verilog)

$RVdst = RVa - RVb$

#### Encoding:

000001 xdddd xaaaa xbbbb xxx xxxx xxxx

#### Usage and Examples:

Subtract two 4x4byte vectors

vv\_sub RV3, RV1, RV2

### 5.3 vv\_mul\_ele

#### Syntax:

vv\_mul\_ele <RVdst>, <RVa>, <RVb>

#### Pseudo Code: (Can be in C/Verilog)

$RVdst = RVa .* RVb$

$RVdst[0] = RVa[0] * RVb[0]$

$RVdst[1] = RVa[1] * RVb[1]$

$RVdst[2] = RVa[2] * RVb[2]$

$RVdst[3] = RVa[3] * RVb[3]$

#### Encoding:

000010 xdddd xaaaa xbbbb xxx xxxx xxxx

#### Usage and Examples:

Elementwise multiply two 4x4byte vectors

vv\_mul\_ele RV3, RV1, RV2

### 5.4 vv\_div

#### Syntax:

vv\_div <RVdst>, <RVa>, <RVb>

**Pseudo Code: (Can be in C/Verilog)**

$$RVdst = RVa / RVb$$

$$RVdst[0] = RVa[0] / RVb[0]$$

$$RVdst[1] = RVa[1] / RVb[1]$$

$$RVdst[2] = RVa[2] / RVb[2]$$

$$RVdst[3] = RVa[3] / RVb[3]$$
**Encoding:**

000011 xdddd xaaaa xbbbb xxx xxxx xxxx

**Usage and Examples:**

Elementwise divide two 4x4byte vectors

vv\_div RV3, RV1, RV2

## 5.5 vf\_add

**Syntax:**

vf\_add <RVdst>, <RVa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

$$RVdst = RVa + Rsb$$

$$RVdst[0] = RVa[0] + Rsb$$

$$RVdst[1] = RVa[1] + Rsb$$

$$RVdst[2] = RVa[2] + Rsb$$

$$RVdst[3] = RVa[3] + Rsb$$
**Encoding:**

001000 xdddd xaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Add 4x4byte vector with a 4byte scaler float

vf\_add RV1, RV2, R1

## 5.6 vf\_sub

**Syntax:**

vf\_sub <RVdst>, <RVa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

$$RVdst = RVa - Rsb$$

$$RVdst[0] = RVa[0] - Rsb$$

$$RVdst[1] = RVa[1] - Rsb$$

$$RVdst[2] = RVa[2] - Rsb$$

$$RVdst[3] = RVa[3] - Rsb$$
**Encoding:**

001001 xdddd xaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Subtract 4byte scaler float from 4x4byte vector

vf\_sub RV1, RV2, R1

### 5.7 vf\_mul

**Syntax:**

vf\_mul <RVdst>, <RVa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RVdst = RVa \* Rsb

RVdst[0] = RVa[0] \* Rsb

RVdst[1] = RVa[1] \* Rsb

RVdst[2] = RVa[2] \* Rsb

RVdst[3] = RVa[3] \* Rsb

**Encoding:**

001010 xdddd xaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Multiply 4x4byte vector with a 4byte scaler float

vf\_mul RV1, RV2, R1

### 5.8 vf\_div

**Syntax:**

vf\_div <RVdst>, <RVa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RVdst = RVa / Rsb

RVdst[0] = RVa[0] / Rsb

RVdst[1] = RVa[1] / Rsb

RVdst[2] = RVa[2] / Rsb

RVdst[3] = RVa[3] / Rsb

**Encoding:**

001011 xdddd xaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Divide 4byte scaler float from 4x4byte vector

vf\_div RV1, RV2, R1

### 5.9 not

**Syntax:**

not <RSd>, <RSs>

**Pseudo Code: (Can be in C/Verilog)**

RSd = ~RSs;

**Encoding:**

001100 ddddd sssss xxxx xxxx xxxx xxxx

**Usage and Examples:**

Inverse the data in the source register in scalar register file and store the result into the destination register.



not R2, R1

Inverse the data in R1 and store the result into R2.

#### 5.10 and

**Syntax:**

and <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa and RSb;

**Encoding:**

001101 ddddd aaaaa bbbbb xxx xxxx xxx

**Usage and Examples:**

The conjunction of two scalar registers and store the result into the destination register.

and R3, R1, R2

*R1 and R2. Store the result into R3.*

#### 5.11 or

**Syntax:**

or <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa or RSb;

**Encoding:**

001110 ddddd aaaaa bbbbb xxx xxxx xxx

**Usage and Examples:**

The disjunction of two scalar registers and store the result into the destination register.

or R3, R1, R2

*R1 or R2. Store the result into R3.*

#### 5.12 xor

**Syntax:**

xor <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa xor RSb;

**Encoding:**

001111 ddddd aaaaa bbbbb xxx xxxx xxx

**Usage and Examples:**

The exclusive disjunction of two scalar registers and store the result into the destination register.

xor R3, R1, R2

*R1 xor R2. Store the result into R3.*

### 5.13 ii\_add

**Syntax:**

ii\_add <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa + RSb;

**Encoding:**

011000 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Add two integers from the scalar register file and store the result into the destination register.

ii\_add RS3, RS1, RS2

Add RS1 and RS2. Then, store the result to RS3.

### 5.14 ii\_sub

**Syntax:**

ii\_sub <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa - RSb;

**Encoding:**

011001 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Subtract one integer from another integer, which both data are from register file. Then, store the result into the destination register.

ii\_sub RS3, RS1, RS2

Subtract RS2 from RS1. Then, store the result to RS3.

### 5.15 ii\_mul

**Syntax:**

ii\_mul <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa \* RSb;

**Encoding:**

011010 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Multiply two integers from the scalar register file and store the result into the destination register.

ii\_mul RS3, RS1, RS2

Multiply RS1 and RS2. Then, store the result to RS3.

### 5.16 ii\_div

**Syntax:**

ii\_div <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa / RSb;

**Encoding:**

011011 dddddd aaaaaa bbbbbb xxx xxxx xxxx

**Usage and Examples:**

Divide one integer by another integer, which both data are from register file. Then, store the result into the destination register.

ii\_div RS3, RS1, RS2

Divide RS1 by RS2. Then, store the result to RS3.

### 5.17 cmp\_i

**Syntax:**

Cmp\_i <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

Flag\_reg = flag(RSa - RSb);

**Encoding:**

011101 xxxxxx aaaaaa bbbbbb xxx xxxx xxxx

**Usage and Examples:**

Compare two integer scalars by computing (RSa - RSb) and output flag to flag register of corresponding thread.

Cmp\_I R1, R2

Compare the integer values in R1 and R2 by computing (R1 - R2) and output flag to flag register of corresponding thread.

### 5.18 ff\_add

**Syntax:**

ff\_add <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa + RSb;

**Encoding:**

010000 dddddd aaaaaa bbbbbb xxx xxxx xxxx

**Usage and Examples:**

Add two floats from the scalar register file and store the result into the destination register.

ff\_add RS3, RS1, RS2

Add RS1 and RS2. Then, store the result to RS3.

### 5.19 ff\_sub

**Syntax:**

ff\_sub <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa - RSb;

**Encoding:**

010001 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Subtract one float from another float, which both data are from register file. Then, store the result into the destination register.

ff\_sub RS3, RS1, RS2

Subtract RS2 from RS1. Then, store the result to RS3.

### 5.20 ff\_mul

**Syntax:**

ff\_mul <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa \* RSb;

**Encoding:**

010010 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Multiply two floats from the scalar register file and store the result into the destination register.

ff\_mul RS3, RS1, RS2

Multiply RS1 and RS2. Then, store the result to RS3.

### 5.21 ff\_div

**Syntax:**

ff\_div <RSd>, <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

RSd = RSa / RSb;

**Encoding:**

010011 ddddd aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Divide one float by another float, which both data are from register file. Then, store the result into the destination register.

ff\_div RS3, RS1, RS2

Divide RS1 by RS2. Then, store the result to RS3.

### 5.22 cmp\_f

**Syntax:**

Cmp\_f <RSa>, <RSb>

**Pseudo Code: (Can be in C/Verilog)**

Flag\_reg = flag(RSa - RSb);

**Encoding:**

010101 xxxxxx aaaaa bbbbb xxx xxxx xxxx

**Usage and Examples:**

Compare two float scalars by computing (RS1 – RSb) and output flag to flag register of corresponding thread.

Cmp\_f R1, R2

Compare the float values in R1 and R2 by computing (R1 - R2) and output flag to flag register of corresponding thread.

### 5.23 s\_write\_high

**Syntax:**

s\_write\_high <RSdst>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

RSdst[31:16] = imm[15:0]

**Encoding:**

100000 ddddd ddddd iiii iiii iiii iiii

**Usage and Examples:**

Fill the highest 2 byte of the scaler register

s\_write\_high R2, #-24

### 5.24 s\_write\_low

**Syntax:**

s\_write\_low <RSdst>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

RSdst[15:0] = imm[15:0]

**Encoding:**

100001 ddddd ddddd iiii iiii iiii iiii

**Usage and Examples:**

Fill the lowest 2 byte of the scaler register

s\_write\_low R2, #24

### 5.25 itof

**Syntax:**

itof <RSd>, <RSs>

**Pseudo Code: (Can be in C/Verilog)**

RSd = float(RSs);

**Encoding:**

10010 ddddd ssss xxxx xxxx xxxx xxxx

**Usage and Examples:**

Convert an integer from the source scalar register to a float and store the result into destination scalar register.

itof R2, R1

Convert the integer in R2 to float and store it into R2.

## 5.26 ftoi

**Syntax:**

itof <RSd>, <RSs>

**Pseudo Code: (Can be in C/Verilog)**

RSd = int(RSs);

**Encoding:**

10011 ddddd ssss xxxx xxxx xxxx xxxx

**Usage and Examples:**

Convert a float from the source scalar register to an integer and store the result into destination scalar register.

itof R2, R1

Convert the float in R2 to integer and store it into R2.

## 5.27 V\_reduce

**Syntax:**

vv\_mul\_dot <RSdst>, <RVsrc>

**Pseudo Code: (Can be in C/Verilog)**

RSdst = RVsrc [0] + RVsrc [1] + RVsrc [2] + RVsrc [3]

**Encoding:**

100100 ddddd xaaaa xxxx xxxx xxxx xxxx

**Usage and Examples:**

Add four element in a vector together

V\_reduce R3, RV1

## 5.28 s\_get\_from\_v

**Syntax:**

s\_get\_from\_v <RSdst>, <RVsrc>, #<index> {0,1,2,3}

**Pseudo Code: (Can be in C/Verilog)**

R2 = RV1[index]

**Encoding:**

100101 ddddd xssss xx xxxx xxxx xxxx ii

**Usage and Examples:**

Get a 4byte scaler from a specified index of a 4x4byte vector

s\_get\_from\_v R2, RV1, #2

### 5.29 S\_sqrt

**Syntax:**

s\_sqrt <RSdst>, <RSsrc>

**Pseudo Code: (Can be in C/Verilog)**

RSdst = sqrt(RSsrc)

**Encoding:**

100110 ddddd sssss xxxx xxxx xxxx xxxx

**Usage and Examples:**

Find square root of a floating point number.

s\_sqrt R2, R1

### 5.30 v\_get\_from\_s

**Syntax:**

v\_get\_from\_s <RVdst>, <RSsrc>, #<index> {0,1,2,3}

**Pseudo Code: (Can be in C/Verilog)**

$RV1[index] = R2$

**Encoding:**

100111 xdddd sssss xx xxxx xxxx xxxx ii

**Usage and Examples:**

Put a 4byte scaler to a specified index of a 4x4byte vector

v\_get\_from\_s RV1, RV2, #2

### 5.31 ii\_addi

**Syntax:**

ii\_addi <RSd>, <RSs>, imm

**Pseudo Code: (Can be in C/Verilog)**

$RSd = RSs + imm$

**Encoding:**

101000 ddddd sssss iiiii iiiii iiiii iiiii

**Usage and Examples:**

Add an integer immediate to an integer from source register in scalar register file. Then, store the result into the destination register.

ii\_addi R2, R1, #7

Add 7 to R1. Then, store the result into R2.

### 5.32 ii\_subi

**Syntax:**

ii\_subi <RSd>, <RSs>, imm

**Pseudo Code: (Can be in C/Verilog)**

$RSd = RSs - imm$

**Encoding:**

101001 ddddd sssss iiiii iiiii iiiii

**Usage and Examples:**

Subtract an integer immediate from an integer from source register in scalar register file. Then, store the result into the destination register.

ii\_subi R2, R1, #7

Subtract 7 from R1. Then, store the result into R2.

### 5.33 ii\_muli

**Syntax:**

ii\_muli <RSd>, <RSs>, imm

**Pseudo Code: (Can be in C/Verilog)**

$RSd = RSs * imm$

**Encoding:**

101010 ddddd sssss iiiii iiiii iiiii

**Usage and Examples:**

Multiply an integer immediate to an integer from source register in scalar register file. Then, store the result into the destination register.

ii\_muli R2, R1, #7

Multiply 7 to R1. Then, store the result into R2.

### 5.34 ii\_divi

**Syntax:**

ii\_divi <RSd>, <RSs>, imm

**Pseudo Code: (Can be in C/Verilog)**

$RSd = RSs / imm$

**Encoding:**

101011 ddddd sssss iiiii iiiii iiiii

**Usage and Examples:**

Divide an integer from the source register in scalar register file by an integer immediate. Then, store the result into the destination register.

ii\_divi R2, R1, #7

Divide R1 by 7. Then, store the result into R2.



### 5.35 s\_load\_4byte

**Syntax:**

s\_load\_4byte <RSdst>, <RSaddress>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

RSdst = mem[RSaddress+imm]

**Encoding:**

101100 ddddd aaaaa iiii iiii iiii iiii

**Usage and Examples:**

Load a 32bit scalar from the address+imm

s\_load\_4byte R2, R1, #7

### 5.36 v\_load\_16byte

**Syntax:**

v\_load\_16byte <RVdst>, <RVaddress>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

RVdst = mem[RVaddress+imm]

**Encoding:**

101111 xdddd aaaaa iiii iiii iiii iiii

**Usage and Examples:**

Load a 4x32bit vector from the address+imm

v\_load\_16byte RV2, R1, #7

### 5.37 s\_store\_4byte

**Syntax:**

s\_store\_4byte <RSsrc>, <RSaddress>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

mem[RSaddress+imm] = RSsrc

**Encoding:**

111100 sssss aaaaa iiii iiii iiii iiii

**Usage and Examples:**

Store a 32bits scalar to the memory address+imm

s\_store\_4byte R2, R1, #7

### 5.38 v\_store\_16byte

**Syntax:**

v\_store\_16byte <RVsrc>, <RVaddress>, #<immediate>

**Pseudo Code: (Can be in C/Verilog)**

mem[RVaddress+imm] = RVsrc

**Encoding:**

111111 xssss aaaaa iiii iiii iiii iiii

**Usage and Examples:**

Store a 4x32bits vector to the memory address+imm

v\_store\_16byte RV2, R1, #7

## 5.39 trace

**Syntax:**

trace

**Pseudo Code: (Can be in C/Verilog)**

Trace: a hardware function call that the argument should be in special registers and returned value will also be in the special registers. Patch Dispatcher will do context switch after calling this function.

Calling: Origin + Direction (retrieve value automatically from RV14 and RV15)

Return: Hit point (X, Y, Z, shader-ID int) + Normal (return automatically to RV14 and RV15)

**Encoding:**

111010 xxxxxx xxxxxx xxxx xxxx xxxx xxxx

**Usage and Examples:**

Trace a ray from its origin along its direction, which are from two different vector registers, to find intersection with triangles.

Trace

## 5.40 Fin

**Syntax:**

FIN

**Encoding:**

111011 xxxxxx xxxxxx xxxx xxxx xxxx xxxx

**Usage and Examples:**

Signal the current program has finished.

## 5.41 be

**Syntax:**

Be LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg == Z)

Program counter <= PC+LABEL;

**Encoding:**

110000 xxxxxx xxxxxx iiii iiii iiii iiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is Z

Bne FUNCTION\_A

#### 5.42 bne

**Syntax:**

Bne LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg != Z)

Program counter <= PC+LABEL;

**Encoding:**

110001 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is not Z.

Bne FUNCTION\_A

#### 5.43 bl

**Syntax:**

Bl LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg == N)

Program counter <= PC+LABEL;

**Encoding:**

110010 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is N, indicating the previous comparison result RSa < RSb.

Bl FUNCTION\_A

#### 5.44 bge

**Syntax:**

Bge LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg == (P || Z))

Program counter <= PC+LABEL

**Encoding:**

110011 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is P or Z, indicating the previous comparison result RSa >= RSb.

Bge FUNCTION\_A

#### 5.45 bg

**Syntax:**

Bg LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg == P)

Program counter <= PC+LABEL;

**Encoding:**

110100 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is P, indicating the previous comparison result RSa > RSb.

Bg FUNCTION\_A

#### 5.46 ble

**Syntax:**

Ble LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog)**

If (flag\_reg == (N || Z))

Program counter <= PC + LABEL;

**Encoding:**

110101 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate if the flag register of corresponding thread is N or Z, indicating the previous comparison result RSa <= RSb.

Ble FUNCTION\_A

#### 5.47 jmp

**Syntax:**

jmp LABEL (#<immediate\_16>)

**Pseudo Code: (Can be in C/Verilog):**

Program counter <= PC+LABEL;

**Encoding:**

110110 xxxxx xxxxx iiiiiiii

**Usage and Examples:**

Jump to the label address represented by the immediate by setting the program counter to it.

Jmp FUNCTION\_A

#### 5.48 return

**Syntax:**

Ret

**Pseudo Code: (Can be in C/Verilog)**

Program counter <= R30;

**Encoding:**

110111 xxxxx 11110 xxxx xxxx xxxx xxxx

### Usage and Examples:

Return subroutine call by setting the program counter to the address saved in R30 by `jmp_link`

## 6 Pseudo Instructions

These instructions can be understood by Assembler, but they are not real instructions that can be execute by the RT core.

| Instructions                                   | Real instructions   |
|--|---|
| <code>S_pop &lt;RSdst&gt;</code>               | <code>s_load_4byte &lt;RSdst&gt; &lt;R29&gt;, ii_addi &lt;R29&gt;, &lt;R29&gt;, -4</code>   |
| <code>S_push &lt;RSsrc&gt;</code>              | <code>s_store_4byte &lt;RSsrc&gt; &lt;R29&gt;, ii_addi &lt;R29&gt;, &lt;R29&gt;, 4</code>   |
| <code>V_pop &lt;RVdst&gt;</code>               | <code>v_load_16byte &lt;RVdst&gt; &lt;R29&gt;, ii_addi &lt;R29&gt;, &lt;R29&gt;, -16</code> |
| <code>V_push &lt;RVsrc&gt;:</code>             | <code>v_store_4byte &lt;RVsrc&gt; &lt;R29&gt;, ii_addi &lt;R29&gt;, &lt;R29&gt;, 16</code>  |
| <code>jmp_link &lt;immediate_16&gt;</code>     | <code>ii_addi R30 R31 4, real_jmp \$l</code>  |
| <code>s_setzero &lt;RSdst&gt;</code>           | <code>xor &lt;RSdst&gt; &lt;RSdst&gt; &lt;RSdst&gt;</code>                                  |
| <code>s_mov &lt;RSdst&gt; &lt;RSsrc&gt;</code> | <code>ii_add &lt;RSdst&gt; &lt;RSsrc&gt; R0</code>  |
| <code>v_mov &lt;RVdst&gt; &lt;RVsrc&gt;</code> | <code>vv_add &lt;RVdst&gt; &lt;RVsrc&gt; R0</code>  |

## 7 Memory Map

From a software point of view, there are two memories: constant and main memory. Other memory addresses are invalid, and they will not return anything.

Constant memory space: 0x20000000 - 0x20007FFF. The constant memory will stay the same no matter the thread ID. The content is generated by software developer and it is the software developers' responsibility to make sure their software and constants agree with each other.

Main memory space: 0x80(Thread ID)0000 - 0x80(Thread ID)FFFF. Each thread will have its own main memory space. We advise software developer to use this memory space as a stack to store local variables and make function calls. At the start of each thread. 0x80(Thread ID)0000 will be placed into R29 automatically so software developer can start push and pop immediately.

Return memory space: 0x80(Thread ID)FFF0 - 0x80(Thread ID)FFFF: Before the thread calling `Fin`, it needs to write the returned value (typically a vector) to this address. Then the hardware will automatically write the value back to the host.

# Micro-architecture

## 1 Overview

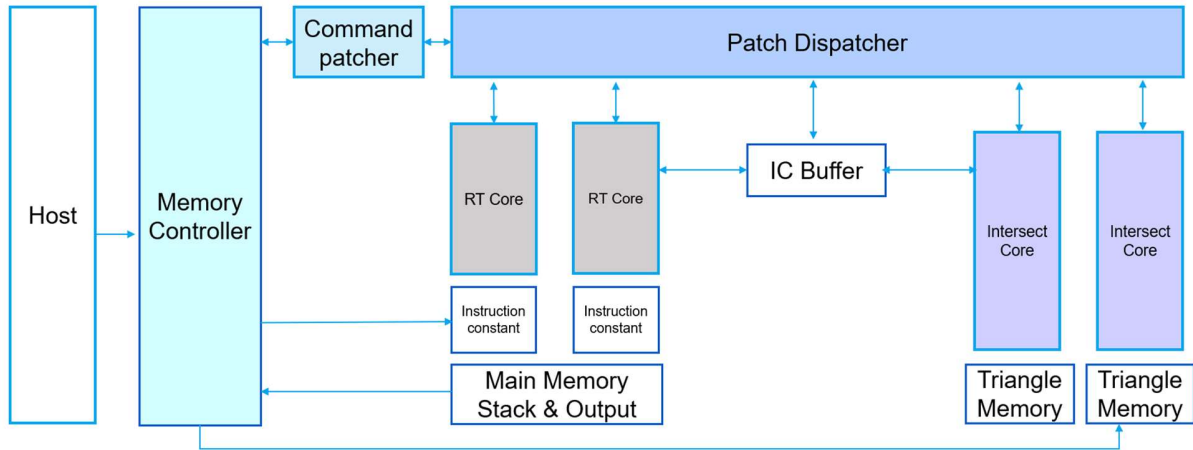


Figure 3 Block diagram of top-level

Our hardware aims to support all the functionality we defined in the Architecture Chapter. Host and self-made memory controller work together through DMA and MMIO interface to move instructions and triangles from host to chip and move the chip to the host.

Command Patcher is a simplified version of the Command Processor we planned. Host will tell the Command Patcher the total number of Pixels and it will break them into a 32-thread chunk. Then it will assign this chunk to Patch Dispatcher. It will also communicate with host when a patch has completed and move the output back to the host.

Patch dispatcher is responsible for scheduling 32 thread to 4 RT Core and 8 IC Core. When a thread evokes trace, it will push this thread to a queue and assign a new thread to this RT Core. It interacts with IC Core in a similar fashion. When all 32 threads have completed, it will signal the Command Patcher that it is ready for the next 32 threads.

RT Core is a generic CPU core with vector and floating-point support and the intersection core is a hardware accelerator. It utilizes the Möller-Trumbore algorithm to achieve a faster execution. The RT Core passes the ray information to IC Buffer. The intersection cores fetch triangle information from triangle memory. A pipeline is used to reduce the down time of the hardware, utilizing the advantages of the algorithm.

Since the Instruction memory, triangle memory and the constant memory is read-only to all the thread, we can simply duplicate them to avoid bank conflict. Each RT Core has its own Instruction and Constant memory, and each IC Core has its own Triangle memory. However, RT Core can read-write main memory therefore we have a unified main memory for all the RT Core.

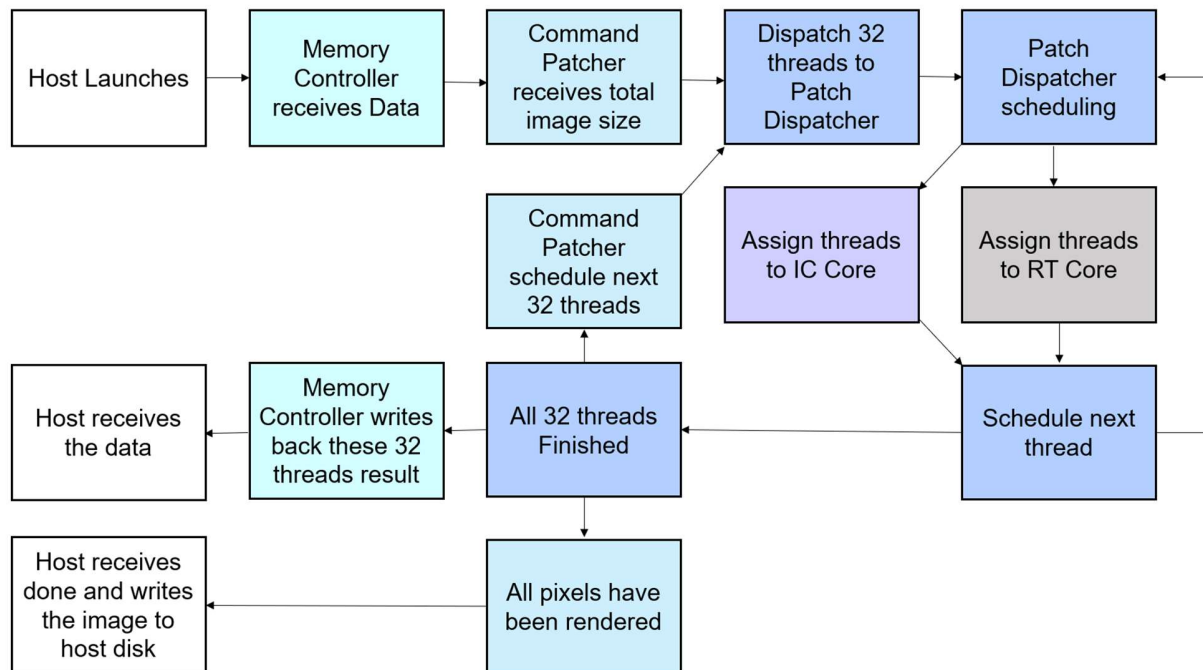


Figure 4 Hardware control flow chart

## 2 Memory Controller

This is the hardware interface for software host. We wrote our own controller to load memory to from the host to our on-chip memory and we did not use the loader provided by the course martials. MMIO passes the starting address, length, and load information of each memory one by one. All information from MMIO is stored into registers. The read controller communicates with the DMA base on the information from MMIO to get content. Then, the data is parsed and formatted for the memory block, and the memory block write-enable signal is raised high. Another purpose of the memory controller is to write output to DMA for software to receive. The write controller requests information from main memory after the ray tracing is done, and the result is stored in registers for an easier write to DMA.

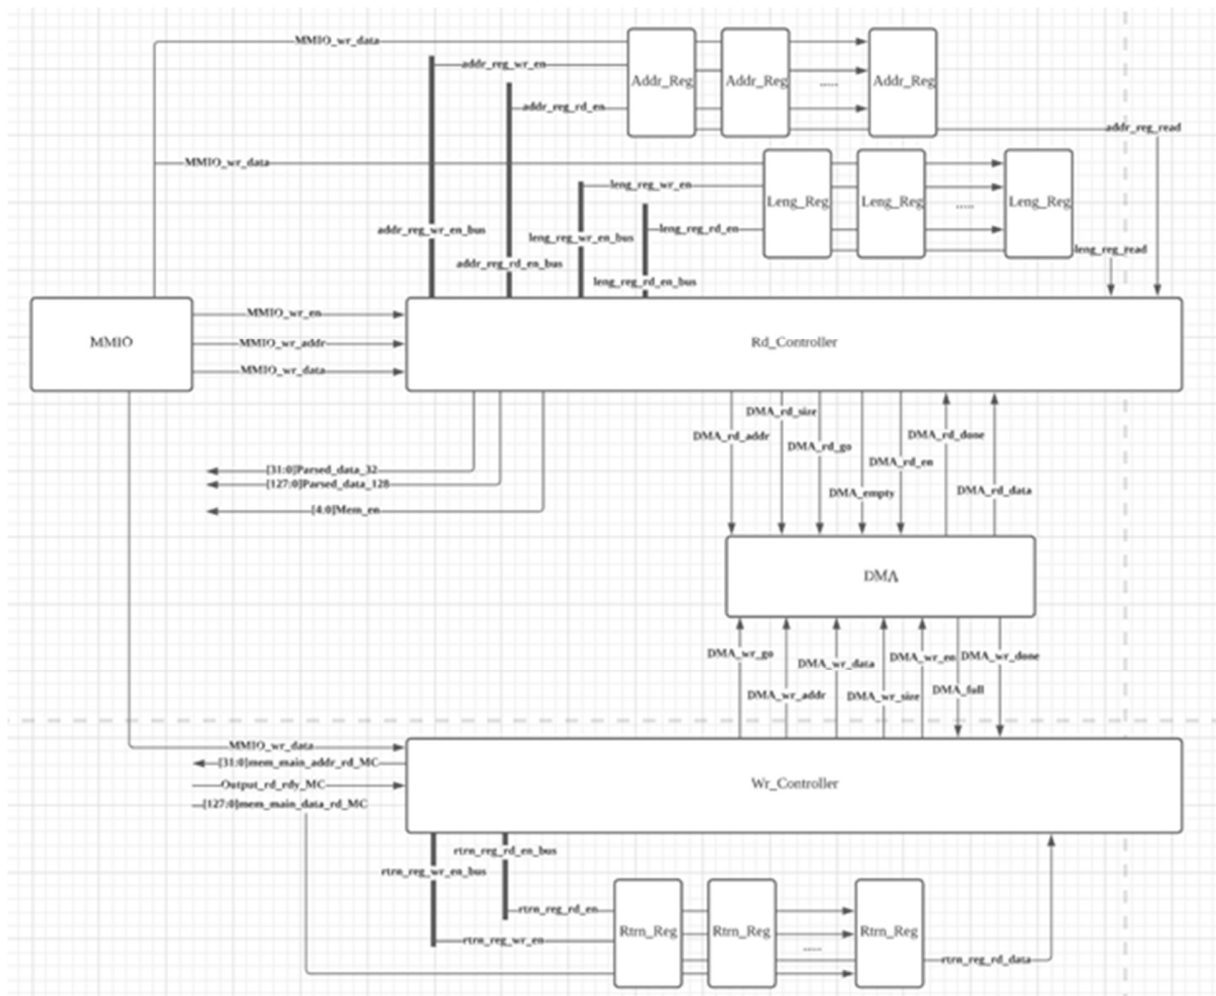


Figure 5 Block Diagram of the internal structure of Memory Controller

#### Global signals

| Signal Name | I/O | Width | Description  |
|-------------|-----|-------|--------------|
| Clk         | I   | 1     | Global Clock |
| Rst_n       | I   | 1     | Global Reset |

#### MC interface

| Signal Name  | I/O | Width | Description  |
|--------------|-----|-------|--|
| MMIO_wr_en   | I   | 1     | Write enable from MMIO indicating data is writing to MC.<br>1 = writing<br>0 = not writing |
| MMIO_wr_addr | I   | 63:0  | The address from software indication what data is passing to MC                            |
| MMIO_wr_data | I   | 63:0  | The data from MMIO   |
| DMA_rd_done  | I   | 1     | Indicating the read from DMA is done.<br>1 = done  |



|                     |   |       |  |
|---------------------|---|-------|--|
|                     |   |       | 0 = reading  |
| DMA_rd_data         | I | 511:0 | The data from DMA  |
| DMA_full            | I | 1     | Indicating the DMA is full or receiving data not yet finished.<br>1 = full<br>0 = not busy |
| DMA_wr_done         | I | 1     | Indicating the writing from DMA is done.<br>1 = done<br>0 = writing                        |
| Output_rd_rdy_MC    | I | 1     | Indicating that the result of ray tracing is ready to be read and write to DMA             |
| Mem_main_data_rd_MC | I | 127:0 | The data from main memory  |
| DMA_empty           | I | 1     | The availability of data from DMA  |
| DMA_rd_addr         | O | 63:0  | The read address of DMA  |
| DMA_rd_size         | O | 9:0   | The size of transfer   |
| DMA_rd_go           | O | 1     | The start signals of reading from DMA  |
| DMA_rd_en           | O | 1     | The receive signal of data from DMA  |
| DMA_wr_go           | O | 1     | The start signals of writing to DMA  |
| DMA_wr_addr         | O | 63:0  | The writing address of DMA   |
| DMA_wr_data         | O | 511:0 | The data write to DMA  |
| DMA_wr_size         | O | 1     | The size of writing data to DMA  |
| DMA_wr_en           | O | 1     | Write enable for DMA   |

Registers.

*Addr\_Reg*

| Depth(bits) | RW | Width (bits) | Function  |
|-------------|----|--------------|---|
| 2:0         | RW | 63:0         | The starting address for content of each memory block |

*Leng\_Reg*

| Depth(bits) | RW | Width (bits) | Function   |
|-------------|----|--------------|--|
| 2:0         | RW | 9:0          | The length of cache line of reading from DMA for each memory block |

*Rtrn\_Reg*

| Depth(bits) | RW | Width (bits) | Function   |
|-------------|----|--------------|--|
| 1:0         | RW | 127:0        | The return data from main memory for writing back to DMA |

### 3 Command Patcher

The Command Patcher initiates a patch of threads for the patch dispatcher. Once the memory controller assigns CP a maximum pixel number and signals the CP to start operating, it calculates the correct pixel

id and writes out the pixel id to patch dispatcher one thread per cycle for 32 cycles. Then it goes back to idle and waits for the done signal from the memory controller again. It also signals the memory controller when it finishes computing all the pixels.

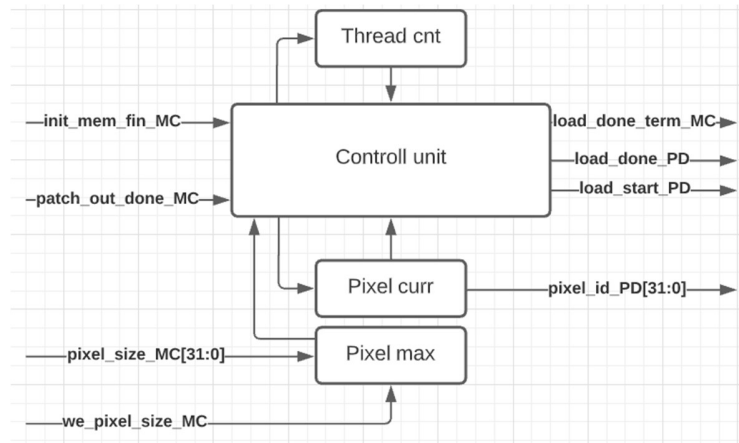


Figure 6 Block Diagram of the internal Structure of the Command Patcher

| Signal Name<br>(name_SRC/TGT) | I/O | Source/Target (blocks) | Description  |
|-------------------------------|-----|------------------------|--|
| Clk                           | I   |                        | Universal clock signal   |
| Rst_n                         | I   |                        | Universal reset low signal   |
| We_pixel_size_MC              | I   | Memory Controller      | Enable high signal for one cycle that loads the maximum pixel id from MC.  |
| [31:0] pixel_size_MC          | I   | Memory Controller      | Maximum pixel size to patch, the processor ends after all pixels have been patched and calculated.                               |
| Init_mem_fin_MC               | I   | Memory Controller      | Done signal from MC indicate that all memories have finished loading, act as the starting signal for CP and the whole processor. |
| Patch_out_done_MC             | I   | Memory Controller      | Done signal from MC indicate that a patch of output has finished transmitting, act as the starting signal for CP.                |
| Load_done_term_MC             | O   | Memory Controller      | Done signal for one cycle to MC indicate that all the pixels have been patched.  |
| Load_start_PD                 | O   | Patch dispatcher       | Starting signal for one cycle that indicate to the PD to start loading new threads, which takes 32 cycles.                       |

|                    |   |                  |   |
|--------------------|---|------------------|---|
| Load_done_PD       | O | Patch dispatcher | Done signal for one cycle after loading of a patch, which indicate to the PD that the job assignment is finished. |
| [31:0] Pixel_id_PD | O | Patch dispatcher | Pixel id assigned for the next thread to PD.  |

## 4 Patch Dispatcher

Patch Dispatcher oversees scheduling 32 threads of operations in between the available number of RT and IC cores. Whenever a RT core reaches the point where the next instruction involves tracing the intersection between the ray and triangles, it would notify the Patch Dispatcher to transfer the process to a dedicated hardware accelerator (IC). RT would send a context switch signal to Patch Dispatcher indicating that it is ready to be transferred. RT would also send its program counter and stack pointer so that this thread could restore its working status when being reassigned.

To re-schedule the thread context switched from RT, the thread ID would be enqueued into a FIFO structure named `rt2ic_queue`. FIFO structure enables a round-robin fashion scheduling so that no thread is being starved. Every clock cycle Patch Dispatcher would assign the FIFO output thread id to an available IC core (if exists, otherwise no operation performed) and dequeue the FIFO. An internal structure is utilized to keep trace of the availability of each RT and IC core and updated when assigning jobs or when the jobs is done. Corresponding program counter and stack pointer are also stored in array type structure indexed by the thread ID.

When IC core finishes tracing and ready to be returned to the RT, it would perform similar procedure by sending a context switch signal to Patch Dispatcher which will enqueue the thread ID into another FIFO structure named `ic2rt_queue`. When any RT core becomes available, Patch Dispatcher will assign the FIFO output thread and dequeue the FIFO. Corresponding program counter, stack pointer and pixel ID are also sent to the RT core to restore the working status of this thread.

Patch Dispatch can deal with multiple number of RT and IC cores. The number of each core are parameterized into the design. For multiple context switch signals from multiple cores at the same clock cycle, Patch Dispatcher would cache the context switch signal and perform the enqueue operation in a fixed descending order once per clock cycle. The number of jobs assigned for each direction (RT to IC or IC to RT) is also one since only one element dequeue is supported by the FIFO structure.

The number of RT core, however, cannot exceed four due to the limitation of the main memory. Main memory supports a maximum of four vector (512bit) memory access bandwidth due to the timing limitation. The number of IC core, on the other hand, could be set to as large as the number of threads.



|                               |   |                          |  |
|-------------------------------|---|--------------------------|--|
| [31:0] pc_RT_in [3:0]         | I | Ray Tracing Core         | Program counter of each RT core.   |
| Intersect_rdy_IC[3:0]         | I | Intersection Core        | Context switch signal ready from IC, indicate that IC is ready to compute a new thread.  |
| [5:0] thread_id_IC_in[3:0]    | I | Intersection Core        | Thread Id indicate which job is running on each IC core.                                 |
| Reload_rdy_CP                 | O | Command Patcher          | Indicate all 64 threads has finished computing and ready to load another 64 new threads. |
| [1:0]Assigned_core_id_RT      | O | Ray Tracing Core         | Core Id indicate which RT core should receive this new job.                              |
| [31:0] pixel_id_RT            | O | Ray Tracing Core         | Pixel id of the new job  |
| [5:0] thread_id_RT_out        | O | Ray Tracing Core         | Thread Id of this new job  |
| [31:0] stack_ptr_RT_out       | O | Ray Tracing Core         | Stack pointer of this new job  |
| [31:0] pc_RT_out              | O | Ray Tracing Core         | Program counter of this new job  |
| [1:0] assigned_core_id_IC     | O | Intersection Core        | Core Id indicate which IC should receive this new job.                                   |
| [5:0] thread_id_IC_out        | O | Intersection Core        | Thread id of this new job.   |
| [3:0] RT2IC_queue_core_id_MIC | O | Intersection Core Memory | RT core id for MIC indicate which RT's output should MIC enqueue.                        |
| [3:0] IC2RT_queue_core_id_MIC | O | Intersection Core Memory | IC id for MIC indicate which IC's output should MIC enqueue.                             |
| RT2IC_queue_en_MIC            | O | Intersection Core Memory | Enable enqueue for RT2IC queue.  |
| IC2RT_queue_en_MIC            | O | Intersection Core Memory | Enable enqueue for IC2RT queue.  |

*Stack pointer/ PC/ Status pool registers*

| Depth | RW | Width | Function   |
|-------|----|-------|--|
| 5:0   | RW | 31:0  | 32-bits Data (stack ptr/PC/status) for each thread is stored in this pool, it can be accessed by its associated thread id. |

## 5 RT core

Our RT CPU adopts a 5-stage pipelined standard design. IF stage will find the next PC counter and send this address location to instruction memory before the next clock edge. DE stage will decode the instructions into multiple control signals. Since our R0 is read only, when an address is a don't care, we always set it to 0. If the current instruction at the DE stage is a branch, it will branch after the Flag writing from Integer ALU.

In EX stage and MEM stage, as most of our operations cannot be done in one cycle. we implemented a state machine to wait for the done signal from Floating point unit or Main memory. Our vectors are always floating so we implement 4 float ALU to compute each element at the same time. One special thing we added was the Vector Reduce Adder. Vectors reduce needs two back-to-back floating points addition and it is too hard to fit in one pipeline stage. We separated this instruction and put it in both EX and MEM stage.

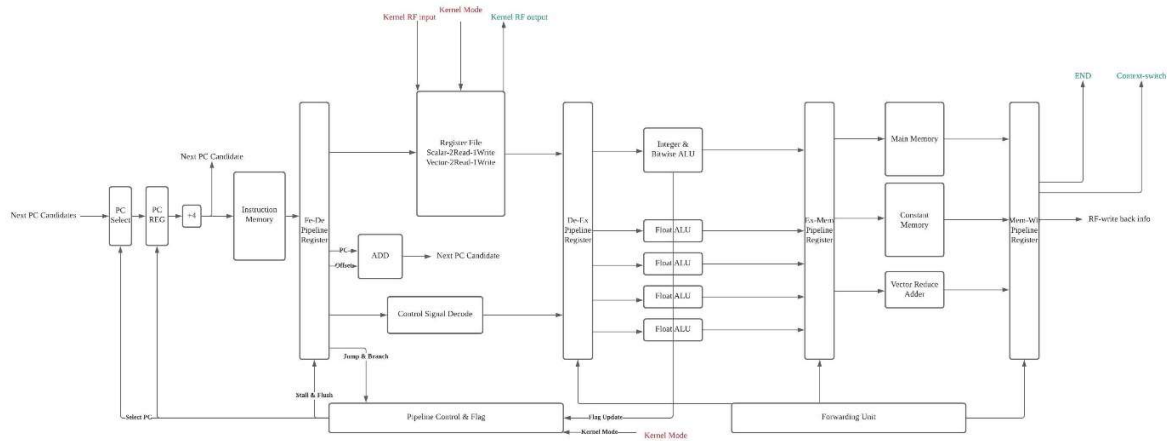


Figure 8 Block Diagram of the internal Structure of the Ray Tracing Core

### Global signals

| Signal Name | I/O | Width | Description  |
|-------------|-----|-------|--------------|
| Clk         | I   | 1     | Global Clock |
| Rst_n       | I   | 1     | Global Reset |

### PD\_RT Interface related signals

| Signal Name | I/O | Width | Description   |
|-------------|-----|-------|---|
| Kernel_mode | I   | 1     | This signal will determine if the core is running on kernel mode. In Kernel mode, the pipeline will stall and the core itself cannot write to anything. The input to the Register file will be switch to signals coming from interface.<br>0 = Normal mode<br>1 = Kernel mode |

|                         |   |       |  |
|-------------------------|---|-------|--|
| PD_scalar_wen           | I | 1     | <b>This signal is useful only in KERNEL mode!</b><br>When enable is active, the register file will write the PD_scalar_wb_data to PD_scalar_wb_address<br>0 = does not write back<br>1 = enable write back |
| PD_scalar_wb_address    | I | 4:0   | <b>This signal is useful only in KERNEL mode!</b><br>Determine which register the data should be write to  |
| PD_scalar_wb_data       | I | 31:0  | <b>This signal is useful only in KERNEL mode!</b><br>Data being write to the Register File   |
| PD_vector_wen           | I | 1     | <b>This signal is useful only in KERNEL mode!</b><br>When enable is active, the register file will write the PD_vector_wb_data to PD_vector_wb_address<br>0 = does not write back<br>1 = enable write back |
| PD_vector_wb_address    | I | 3:0   | <b>This signal is useful only in KERNEL mode!</b><br>Determine which register the data should be write to  |
| PD_vector_wb_data       | I | 31:0  | <b>This signal is useful only in KERNEL mode!</b><br>Data being write to the Register File   |
| PD_scalar_read_address1 | I | 4:0   | <b>This signal is useful only in KERNEL mode!</b><br>The register file will output the information in the register specified by this address. Data will be ready at port 1                                 |
| PD_scalar_read_address2 | I | 4:0   | <b>This signal is useful only in KERNEL mode!</b><br>The register file will output the information in the register specified by this address. Data will be ready at port 2                                 |
| PD_vector_read_address1 | I | 3:0   | <b>This signal is useful only in KERNEL mode!</b><br>The register file will output the information in the register specified by this address. Data will be ready at port 1                                 |
| PD_vector_read_address2 | I | 3:0   | <b>This signal is useful only in KERNEL mode!</b><br>The register file will output the information in the register specified by this address. Data will be ready at port 2                                 |
| PD_scalar_read1         | O | 31:0  | <b>This signal is useful only in KERNEL mode!</b><br>Supply data to specified by a PD_scalar_read_address1   |
| PD_scalar_read2         | O | 31:0  | <b>This signal is useful only in KERNEL mode!</b><br>Read data specified by a PD_scalar_read_address2  |
| PD_vector_read1         | O | 127:0 | <b>This signal is useful only in KERNEL mode!</b><br>Read data specified by a PD_vector_read_address1  |
| PD_vector_read2         | O | 127:0 | <b>This signal is useful only in KERNEL mode!</b><br>Read data specified by a PD_vector_read_address2  |

|                |   |   |   |
|----------------|---|---|---|
| END            | O | 1 | This signal will notify the Patch dispatcher that the current thread running on the core has finished. Then patch dispatcher can assign other thread to this core.  |
| Context_switch | O | 1 | This signal will notify the interface that this thread wants to do a ray trace and this thread handled all register other than Stack based Pointer and PC. Then Interface will read tracing data, stack pointer and PC from the core then send to Patch dispatcher and IC memory. |

### *Memory Related Signals*

| Signal Name    | I/O | Width | Description  |
|----------------|-----|-------|--|
| MRTI_addr      | O   | 31:0  | The address of next PC instruction                   |
| MRTI_data      | I   | 31:0  | The current instruction read from instruction memory |
| MEM_addr       | O   | 31:0  | The address of memory address                        |
| MEM_data_write | O   | 127:0 | Memory data write to the memory                      |
| MEM_data_read  | I   | 127:0 | Memory data read from the memory                     |
| MEM_read_en    | O   | 1     | Enable memory read                                   |
| MEM_write_en   | O   | 1     | Enable memory write                                  |
| MEM_done       | I   | 1     | Indicate the read data is ready                      |
| MEM_s_or_v     | O   | 1     | Indicate the current operation is scalar or vector   |

## 6 IC buffer

The IC buffer is the storage space between intersection cores and ray-tracing cores. Once any core has finished its calculations, the patch dispatcher will control the buffer to select and load the corresponding data into the buffer queue. After the same thread has been assigned again on the other type of core. IC cores and RT cores then read from this memory and dequeue the buffer when starting. This process utilizes two queues which support random-access writes to store and pass data. All of the input data will be written into the tail of the queue. All IC cores and RT cores connect to the data input and MUXs are used to choose input based on PD signals. For example, once an IC has finished and signals the PD to push data into the queue, it waits until it has been re-assigned to a RT core, and then dequeues the data from the front of the queue. Since enqueue and dequeue are carefully managed by PD, the data remains synchronized with the PD assignment.



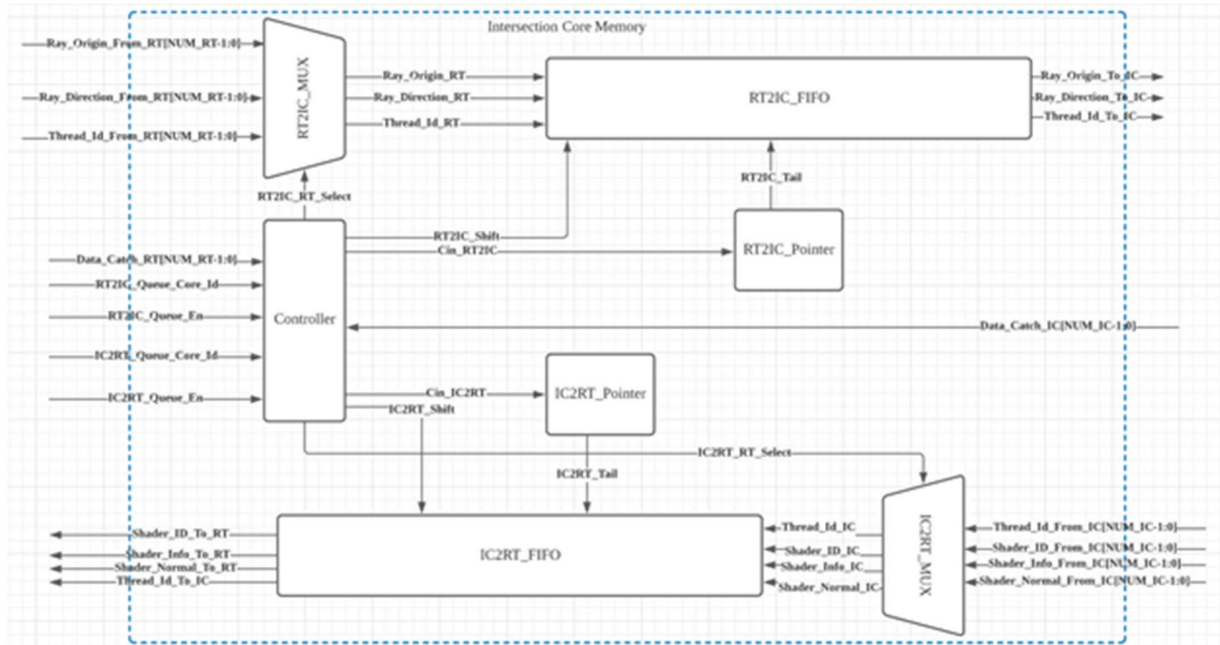


Figure 9 Block Diagram of the internal Structure of intersection Core buffer

#### Global signals

| Signal Name | I/O | Width | Description  |
|-------------|-----|-------|--------------|
| Clk         | I   | 1     | Global Clock |
| Rst_n       | I   | 1     | Global Reset |

#### ICM interface

| Signal Name            | I/O | Width | Description  |
|------------------------|-----|-------|--|
| Ray_origian_RT[3:0]    | I   | 95:0  | The ray origin transferred from RT cores.  |
| Ray_direction_RT[3:0]  | I   | 95:0  | The direction of the ray transferred from RT cores.  |
| Dequeue_RT             | I   | 1     | Dequeue signal from RT indicate that RT has finished reading the data and it is safe to discard the data in queue. |
| Q_en_rt2ic_PD          | I   | 1     | Enqueue signal from PD indicates that RT has finished calculation and should enqueue data from the selected RT.    |
| Core_id_rt2ic_PD [3:0] | I   | 1     | One hot selects signal indicates which RT core to read data from.  |
| Q_en_ic2rt_PD          | I   | 1     | Enqueue signal from PD indicates that IC has finished calculation and should enqueue data from the selected IC.    |
| Core_id_ic2rt_PD [3:0] | I   | 1     | One hot selects signal indicates which IC core to read data from.  |

|                     |   |       |  |
|---------------------|---|-------|--|
| Shader_Info_IC[3:0] | I | 127:0 | The coordinate and shader id of intersection points transferred from IC cores.                                     |
| Normal_IC[3:0]      | I | 95:0  | The normal vector of the triangle on the intersection points transferred from IC cores.                            |
| Dequeue_IC          | I | 1     | Dequeue signal from IC indicate that IC has finished reading the data and it is safe to discard the data in queue. |
| Ray_origian_IC      | O | 95:0  | The ray origin data output for IC cores.   |
| Ray_direction_IC    | O | 95:0  | The direction of the ray output for IC cores.  |
| Shader_Info_RT      | O | 127:0 | The coordinate of intersection points output for RT cores.   |
| Shader_Normal_RT    | O | 95:0  | The normal vector of the triangle on the intersection points output for RT cores                                   |

#### *RT2IC Queue*

| Depth(bits) | RW | Width (bits) | Function  |
|-------------|----|--------------|---|
| 31:0        | RW | 191:0        | A queue with the supporting of random access write to the tail of the queue for storing ray information from RT cores and pass it to IC cores. A pointer points to the tail for writing new data. |

#### *IC2RT Queue*

| Depth(bits) | RW | Width (bits) | Function  |
|-------------|----|--------------|---|
| 31:0        | RW | 223:0        | A queue with the supporting of random access write to the tail of the queue for storing ray information from RT cores and pass it to IC cores. A pointer points to the tail for writing new data. |

## 7 IC core

Our intersection core (IC) is implemented with a 4-stage pipelined design. The IC utilizes Möller-Trumbore algorithm for fast detection of intersections. Floating point addition, subtraction, multilocation, reciprocal, and comparison are used in this module.

Based on the nature of the algorithm, the process can be divided in to 4 serial stages, allowing us to pipeline our detections. Furthermore, stage pipelines can be switched to using multiple FIFOs with different lengths to meet the latency between input signals. By doing this, the throughput can be further increased to one cycle per output. However, this “FIFO design” has two downsides. First, it requires a fixed clk frequency since the latency of the IPs are depended on the clk frequency. The FIFOs cannot be too long nor too short. It would lead to false results. Secondly, the bottle neck of the performance would then be fetching from Triangle Memory. It requires at least 6 cycles to fetch. Therefore, using a stage design can best balance the throughput and robustness of the design.

The detection goes through all of the triangles and finds the closest intersecting triangle. The IC receives the starting signal from PD and ray information from IC memory. The controller would fetch the triangle data with triangle memory. Based on the return data, it can be determined where the ray intersects the triangle, if it intersects the triangle at all. Finally, the IC core returns the necessary data back to IC memory. If there is no intersection with any triangle, then shader ID 0 is returned to indicate. We further notice that there is one more condition for intersecting with a triangle that is not present in the algorithm, which is the distance  $t$  between origin of the light and the triangle should be positive. The original algorithm only comparing the determinant,  $u$ , and  $v$  to determine hitting or not. This is crucial for getting the right shadow of the final rendering.

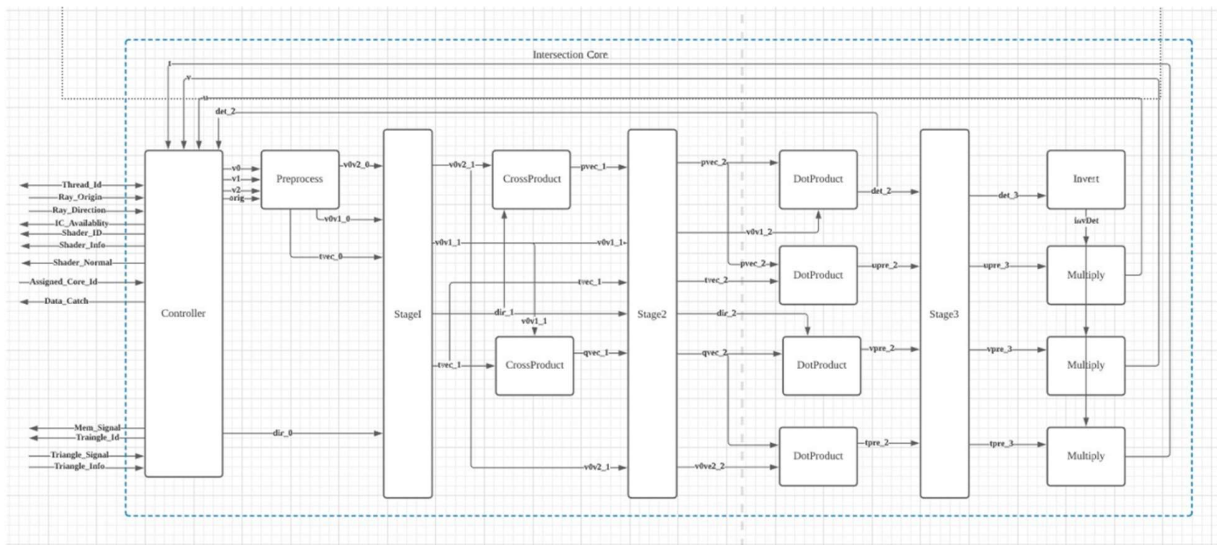


Figure 10 Block Diagram of the internal Structure of the Intersection Core

#### Global signals

| Signal Name | I/O | Width | Description  |
|-------------|-----|-------|--------------|
| Clk         | I   | 1     | Global Clock |
| Rst_n       | I   | 1     | Global Reset |

#### IC interface

| Signal Name      | I/O | Width | Description   |
|------------------|-----|-------|---|
| Thread_Id        | I/O | 5:0   | Documenting the thread id of the intersection detection. This is passed from and sent back to IC memory       |
| Ray_origin       | I   | 127:0 | The origin of the ray   |
| Ray_direction    | I   | 127:0 | The direction of the ray  |
| Assigned_core_Id | I   | 1:0   | The start signals from PD. Only receiving ray data from IC memory when this signal matches with the core ID   |
| Triangle_Signal  | I   | 1:0   | The triangle information ready signal. Only receiving triangle data when this signal matches with the core ID |

|                 |   |       |   |
|-----------------|---|-------|---|
| Triangle_Info   | I | 127:0 | The triangle information receiving from triangle memory   |
| IC_availability | O | 1     | Indicating that the IC core finished its job and is available.<br>1 = available<br>0 = busy                                       |
| Shader_Id       | O | 8:0   | Return the ID of closest intersecting triangle. If no intersection, return 0  |
| Shader_Info     | O | 127:0 | The coordinate of intersection point. If no intersection, don't care  |
| Shader_Normal   | O | 127:0 | The normal vector of the triangle on the intersection point. If no intersection, don't care                                       |
| Data_catch      | O | 1     | Indicating that the information from IC memory is caught. Stay high for 1 cycle after caught.<br>1 = caught<br>0 = not yet caught |
| Mem_signal      | O | 1     | Indicating that the core is requesting triangle information.<br>1 = requesting<br>0 = not requesting                              |
| Traingle_Id     | O | 8:0   | The id of requesting triangle   |

## 8 IP cores

Ray tracing has many calculations in the algorithms, including both integers and floating point. In order to have a better performance per resource, we decided to use IP cores for the fundamental calculations for the integers and floating points.

The IP cores are Intel Arithmetic IPs for fixed point and floating point. By setting the fraction of the fixed point to 0, we can use it as integers. Even though Intel IP cores are very powerful, it has some inconveniences for us to integrate them into our design. First, almost all the IPs need a couple cycles to calculate. However, there is no 'done' signal for us to now whether the calculation has finished or not. Luckily, the IP provided a estimated latency for different clock frequency. We could use those information to write a counter to generate a 'done' signal for outer level. Secondly, as mentioned above, the latency for different IPs is different under different clock frequency. And, this information is not documented in anywhere since it is device specific. We have to record each latency manually while generating the IPs. Third, the simulation of the IPs is very difficult. All of them requires a library called 'fourteenm' which is not included for the standard version Modelsim. After many trails and errors, we found out that the Pro version of Modelsim is needed to simulate IPs. Even though there is a free starter license for the Pro version, it limits the line count of the simulated project to 10,000 lines. When we tested our ALUs, we found out that the IP themselves had already surpassed the line limitation.

Even though there is many hardships with using the Intel IPs, the performance is promising. The best benefit of them all is that it guarantees the timing for synthesis.

## 9 Memory

Memory space is separated into several groups on the top level to serve specific purposes. Specifications of each memory module are demonstrated in the following table. We have optimized the memory structure to best support our multicore design structure.

### 9.1 Memory overview

| Name                        | Type             | Internal Data Width (bit) | IO Data Width (bit) | Depth | Bank # | Size (bit)         |
|-----------------------------|------------------|---------------------------|---------------------|-------|--------|--------------------|
| Command Patcher Instruction | Simple Dual Port | 32                        | 32                  | 512   | 1      | 16 K               |
| RT Instruction x 4          | Simple Dual Port | 32                        | 32                  | 4096  | 1      | 128 x 4 = 512 K    |
| RT Constant x 4             | Simple Dual Port | 32                        | 32                  | 512   | 1      | 16 x 4 = 64 K      |
| Triangle x 4                | Simple Dual Port | 32                        | 32                  | 2048  | 4      | 16 x 4 x 4 = 256 K |
| Main Memory                 | Simple Dual Port | 32                        | 32 x 4 x 4          | 4096  | 256    | 32 M               |

Figure 11 Memory Specification table

### 9.2 Instruction memory

RT instruction memory module stores the software instructions that will be run on the RT core. The instructions will be initially loaded from the memory controller incrementally and terminate with a control signal assigned by MC. There is a single ram bank within the module to support one read from RT or one write from MC each cycle.

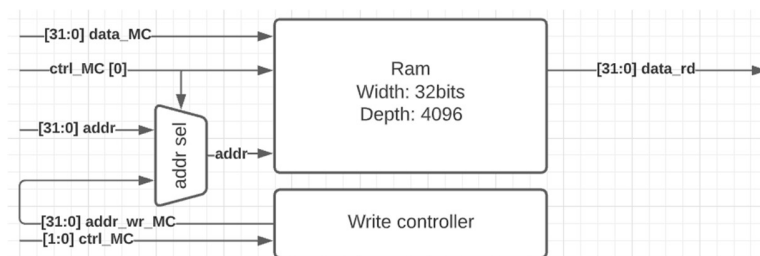


Figure 12 Block diagram of the internal structure of the Ray Tracing instruction memory

| Signal Name    | I/O | Source/Target<br>(blocks) | Description   |
|----------------|-----|---------------------------|---|
| clk            | I   |                           | Universal Clock signal  |
| Rst_n          | I   |                           | Universal reset signal  |
| [1:0] ctrl_MC  | I   | Memory controller         | Control signal for memory controller.<br>Bit 0 is the write enable signal.<br>Bit 1 is the write finished signal. |
| [31:0] data_MC | I   | Memory controller         | Data to write to the address by MC.   |
| [31:0] addr    | I   | Ray Tracing               | Address to read from by RT.   |
| [31:0] data_rd | O   | Ray Tracing               | Data read from the address by RT.   |

#### Memory Bank

| Bits | RW | Width | Function  |
|------|----|-------|---|
| 11:0 | RW | 31:0  | Memory bank used to store the instructions of RT. |

### 9.3 Constant memory

The RT constant memory module stores the software constants that are needed for running the RT core. The constants will be initially loaded from the memory controller incrementally, terminating with a control signal assigned by MC. There is a single ram bank within the module to support one read from RT or one write from MC each cycle.




Figure 13 Block diagram of the internal structure of the Ray Tracing constant memory

| Signal Name   | I/O | Source/Target<br>(blocks) | Description   |
|---------------|-----|---------------------------|---|
| clk           | I   |                           | Universal Clock signal  |
| Rst_n         | I   |                           | Universal reset signal  |
| [1:0] ctrl_MC | I   | Memory controller         | Control signal for memory controller.<br>Bit 0 is the write enable signal.<br>Bit 1 is the write finished signal. |

|               |   |                   |                                     |
|---------------|---|-------------------|-------------------------------------|
| [31:0]data_MC | I | Memory controller | Data to write to the address by MC. |
| [31:0]addr    | I | Ray Tracing       | Address to read from by RT.         |
| [31:0]data_rd | O | Ray Tracing       | Data read from the address by RT.   |

#### Memory Bank

| Bits | RW | Width(bits) | Function                                      |
|------|----|-------------|---|
| 8:0  | RW | 32          | Memory bank used to store the constant of RT. |

### 9.4 Main memory

The RT Main memory is a unified memory space accessible to software, which is reserved for the stack space of each thread and ray tracing output. The RT main memory is word-addressable and supports four parallel vectors (128 bits width) reads/writes at the same time. The internal structure of this memory module consists of 128 banks of RAM with 4 banks per thread to minimize the memory conflict and maximize throughput. Since the logic essentially requires a mapping from 4x4x32 bits to 128x32 bits, in order to optimize for timing and throughput, we have developed a pipelined version of this memory. We first have made a mapping from 4x4x32 to 8x4x32 and then a second mapping from 8x4x32 to 128x32 in a stage of the pipeline to limit the size of the selecting logic. Thus, with the optimization, we are able to have the maximum throughput of 512 bits per cycle and thus vastly improved our timing as well as throughput.

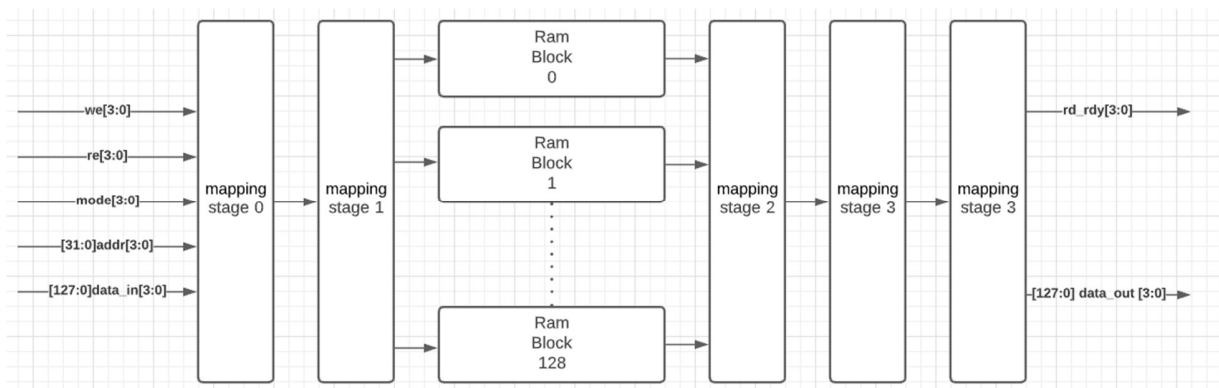


Figure 14 Block diagram of the Structure of the Main memory (Note: the mapping logics are hidden)

| Signal Name | I/O | Source/Target (blocks)          | Description                       |
|-------------|-----|---------------------------------|-----------------------------------|
| clk         | I   |                                 | Universal Clock signal            |
| rst_n       | I   |                                 | Universal reset signal            |
| we [3:0]    | I   | Ray tracing / Memory controller | Write enable signal for each port |

|                        |   |  |  |
|------------------------|---|--|--|
| re [3:0]               | I | Ray tracing /<br>Memory controller       | Read enable signal for each port   |
| mode [3:0]             | I | Ray tracing                              | Mode signal for different mode of write for<br>RT core<br>0: Scaler write<br>1: Vector write     |
| [31:0] addr [3:0]      | I | Ray tracing /<br>Memory controller       | Read/write address to read data from main<br>memory  |
| [127:0] data_in [3:0]  | I | Ray tracing /<br>Memory controller       | Vector Data input written into the main<br>memory  |
| rd_rdy [3:0]           | O | Ray tracing /<br>Memory controller       | Ready signal to indicate the read operation is<br>complete and the data is available at data_out |
| [127:0] data_out [3:0] | O | Memory controller /<br>Memory controller | Vector Data read from the main memory<br>requested by the memory controller                      |

#### *Ram Bank[127:0]*

| Depth | RW | Width | Function  |
|-------|----|-------|---|
| 11:0  | RW | 31:0  | Memory bank on which the unified main memory is mapped to<br>enable parallel read / write without conflict. Each 4 linear Ram banks<br>will be combined into a 2D structure on which the unified memory<br>address maps in a row major fashion for a single thread, and total<br>32x4 banks for 32 threads. |

### 9.5 Triangle memory

The triangle memory module stores the vertices of all triangles, their corresponding x, y, z coordinate data, and their shader ids which are used in the software for image calculation. Due to the overlapping of vertices among triangles, it would be wasteful to store all information on each vertex of each triangle in memory. Instead, we have designed a special data structure to optimize space utilization. Triangle memory module stores information on each individual vertex directly in the index memory bank, and stores a mapping table in the vertex memory banks. This is to find the corresponding vertex data of any given triangles using the mapping table. As a result, it reduces the memory required to store the triangles by a factor up to 3. On the other hand, we have also optimized the design to improve the throughput. To read the 320 bits data for one triangle (3x32 bits per-vertex and 32 bits shader id), the process takes up to 5 cycles, thus a prefetch mechanism is implemented to reduce latency. When the memory is idle from IC's quest, it starts prefetching the next triangle's data. Since IC is comparatively slower, prefetching lowers the latency to at minimum 1 cycle, thus increases relative throughput to 320 bits per cycle.



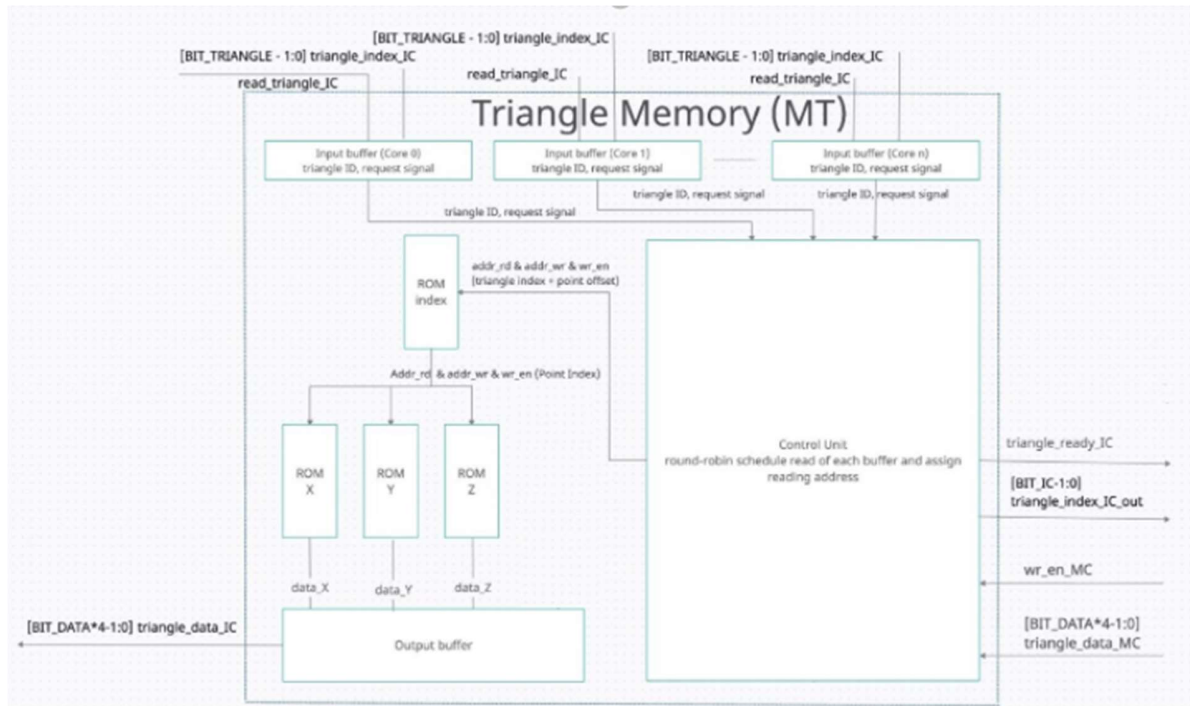


Figure 15 Block Diagram of the internal structure of Triangle Memory

| Signal Name          | I/O | Source/Target (blocks) | Description  |
|----------------------|-----|------------------------|--|
| clk                  | I   |                        | Universal Clock signal   |
| rst_n                | I   |                        | Universal reset signal   |
| re_IC                | I   | Intersection core      | Signal indicate IC want to read triangle data from the triangle memory.  |
| [8:0] triangle_id_IC | I   | Intersection core      | Id for triangle indicate which triangle IC want to read.   |
| [127:0] data_MC      | I   | Memory controller      | Triangle data to write in during MC memory write.  |
| We_MC                | I   | Memory controller      | Write enable signal for MC data write.   |
| Done_MC              | I   | Memory controller      | Terminate signal from MC indicate that the loading stage of the memory is completed.   |
| rdy_IC               | O   | Intersection core      | Signals indicate the triangle read is ready.   |
| Not_valid_IC         | O   | Intersection core      | Terminate signal for IC indicate that the address id read is not a valid triangle id which acts as the termination signal for IC's intersection calculation. |

|                   |   |                   |   |
|-------------------|---|-------------------|---|
| Rdy_MC            | O | Memory controller | Singal indicate the triangle write from MC is done.                 |
| [95:0] vertex1_IC | O | Intersection core | X, Y, Z data information for the first vertex of the triangle read  |
| [95:0] vertex2_IC | O | Intersection core | X, Y, Z data information for the second vertex of the triangle read |
| [95:0] vertex3_IC | O | Intersection core | X, Y, Z data information for the third vertex of the triangle read  |
| [95:0] sid_IC     | O | Intersection core | Shader id data for the triangle read, used in software calculation. |

#### Memory Bank Index

| Bits | RW | Function  |
|------|----|---|
| 10:0 | RW | Memory bank used to store triangle data, this bank is used to store the vertex id and shader id of all triangles, each triangle takes 4 memory entries. |

#### Memory Bank X/Y/Z

| Depth | RW | Width | Function  |
|-------|----|-------|---|
| 10:0  | RW | 31:0  | Memory bank used to store X / Y / Z coordinates and alpha value of each vertex. |

#### Output Buffer

| Bits  | RW | Function  |
|-------|----|---|
| 320:0 | RW | Output buffer holds all the data output from triangle memory loaded |

## 10 Coding standard

System Verilog:

All wire and interface name are in lowercase. Use underscore for separating words.

Name wire in the format: wirename\_SS\_DD.

Example: thread\_id\_cp\_pd

Name interface in the format of: interfacename\_XX

Example: thread\_id\_pd

If input and output have the same name, add \_in or \_out

Example: thread\_id\_pd\_in or thread\_id\_pd\_out

| Module name       |    | Wire name |      |
|-------------------|----|-----------|------|
| Command Patcher   | CP | Pointer   | ptr  |
| Patch Dispatcher  | PD | Ready     | rdy  |
| Ray Trace Core    | RT | Address   | addr |
| Intersection Core | IC | Read      | rd   |
| Triangle Memory   | MT | Write     | wr   |

|                          |      |             |       |
|--------------------------|------|-------------|-------|
| IC Memory                | MIC  | Enable      | en    |
| CP Memory                | MCP  | Instruction | inst  |
| RT Main Memory           | MRTM | Constant    | const |
| RT instruction<br>memory | MRTI | Branch      | br    |
| RT constant memory       | MRTC |             |       |
| Memory Controller        | MC   |             |       |